

MuJoCoXR

Linking MuJoCo and a VR environment

2024-07-22

Youenn Le Jeune
INSA Rennes / DIAG Sapienza



I. Introduction and State of the Art

The MuJoCo physics engine is a great tool that contains many useful features that are nowhere to be found in other engines. Such features are, for instance, realistic soft body simulation. Making interactive simulations is possible through the provided API, in C or in Python. It is easy to integrate motion-capture devices and to get feedback from sensors.

However, the rendering pipeline can be difficult to comprehend. There are multiple pre-made visualization solutions: an interactive one for your MJCF files via the “Simulate”¹ code sample, a passive one that you can use in your Python application to show your interactive simulation in real-time², a full Renderer Python class³ that can be really useful for notebooks and, if nothing else is suitable, API functions to render directly in an OpenGL context⁴.

As we can see, there is no built-in way to visualize a MuJoCo simulation in a Virtual Reality environment. At the time of writing, some work is being done to add VR support to the “Simulate” application⁵. Nevertheless, it is still not the best choice if you want to fully immerse people in your simulation⁶. One possible solution is to use the Unity Plug-in for MuJoCo⁷ and configure the Unity application to display in VR (*not tested*), but again that is not ideal: we do not necessarily need a whole game engine to run a simple simulation. Thus, we will program our own solution to display a MuJoCo simulation in a VR headset.

Regarding the VR part, there are multiple solutions to render to a headset:

- Directly use the vendor-specific API (Meta Quest, HTC Vive...). This requires code to be remade for each new device we want to support.
- Use OpenVR, which contains support for VR headsets from multiple vendors but it is tied to Steam.⁸
- Use OpenXR, which is a free standard not tied to any VR company and implemented by all major VR headsets on the market.⁹

We have chosen to use OpenXR because it seems to be the standard that will be mainly used in the future. There are Python bindings available¹⁰. It supports multiple graphics API including OpenGL that we will use because it is what MuJoCo can render to.

To sum up, our solution will render MuJoCo on a VR headset using OpenGL via the OpenXR standard, all of that in Python.

¹<https://mujoco.readthedocs.io/en/stable/programming/samples.html#sasimulate>

²<https://mujoco.readthedocs.io/en/stable/python.html#passive-viewer>

³<https://github.com/google-deeppmind/mujoco/blob/main/python/mujoco/renderer.py>

⁴<https://mujoco.readthedocs.io/en/stable/programming/visualization.html>

⁵<https://github.com/google-deeppmind/mujoco/pull/1452>

⁶The “Simulate” application contains a lot of UI elements to control the simulation and visualization, which breaks immersion.

⁷<https://mujoco.readthedocs.io/en/stable/unity.html>

⁸<https://github.com/ValveSoftware/openvr>

⁹<https://www.khronos.org/openxr/>

¹⁰<https://github.com/cmbruns/pyopenxr>

II. Theory

II - 1. Graphics 101

To render something on a screen or a VR headset, computers use graphic cards (GPUs). Those graphic cards receive orders through Graphic APIs such as OpenGL, Direct3D or more recently, Vulkan. Each GPU supports different versions of those APIs, and old GPUs do not even support some APIs.

A Graphic API consists of a large set of instructions related to graphics: draw a line from here to there, clear the screen with this color, draw this texture. OpenGL instructions are all prefixed with `gl` and make heavy use of constants. For instance, a possible instruction is `glClear(GL_COLOR_BUFFER_BIT)`.

Graphic APIs instructions are not only used to *draw*: with the support of *framebuffers*, it is possible for instance to draw on some in-memory texture and then read it to draw on top of another buffer with a different scale. Framebuffers are a “collection” of renderbuffers and textures and can have multiple *attachements*: colors, depth and stencil.

OpenGL is built on the principle of *extensions*.

Graphic APIs work with a *context*: to use their functions, a context must be bound to the thread. It contains references to all GL objects created within it. A context is usually tied to a window. To create those contexts, we usually use dedicated libraries such as GLFW which allows to create a window and attach the associated context to the calling thread.

In Python, there exist a binding for OpenGL: `pyopengl`¹¹.

II - 2. OpenXR

OpenXR is a standard implemented by pretty much all VR devices. It provides methods for most of the features: displaying images, getting head position in the room, getting controller positions, rendering haptic feedback, enable passthrough for compatible headsets, and so on. For vendor-specific features, extensions are present in OpenXR to use them.

To render images to the eyes, OpenXR uses the concept of *swapchains*. A swapchain is a collection of framebuffers that display sequentially to a screen. It allows to draw on a “back” framebuffer while another “front” one is being displayed on the screen.¹²

In this project, we will make use of only one “stereo” swapchain, which will contain one image for both eyes at the same time. It is also possible to have multiple swapchains, e.g. one per eye.

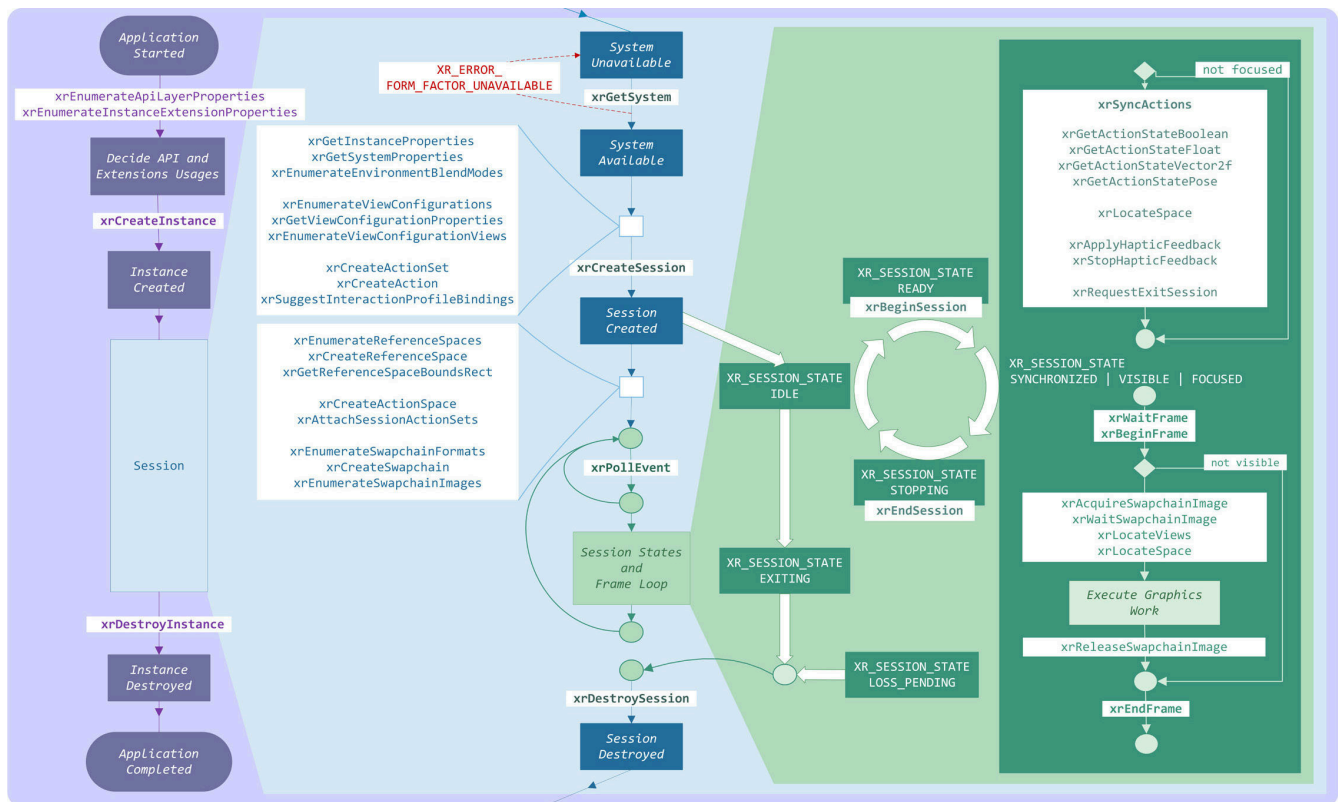
In order to create an OpenXR-compatible application, a precise suite of operations must be followed (see Figure 1 for details):

1. Available extensions are fetched to see if the ones needed are present (for instance, the extension that tells OpenXR to use the OpenGL graphics API, the debug utils extension...)
2. An “instance” is created with the application informations and extensions list. Future methods will use this instance.

¹¹<https://pyopengl.sourceforge.net/>

¹²<https://raphlinus.github.io/ui/graphics/gpu/2021/10/22/swapchain-frame-pacing.html>

Figure 1 — Lifecycle of an OpenXR application¹³
Not all functions are necessarily used.

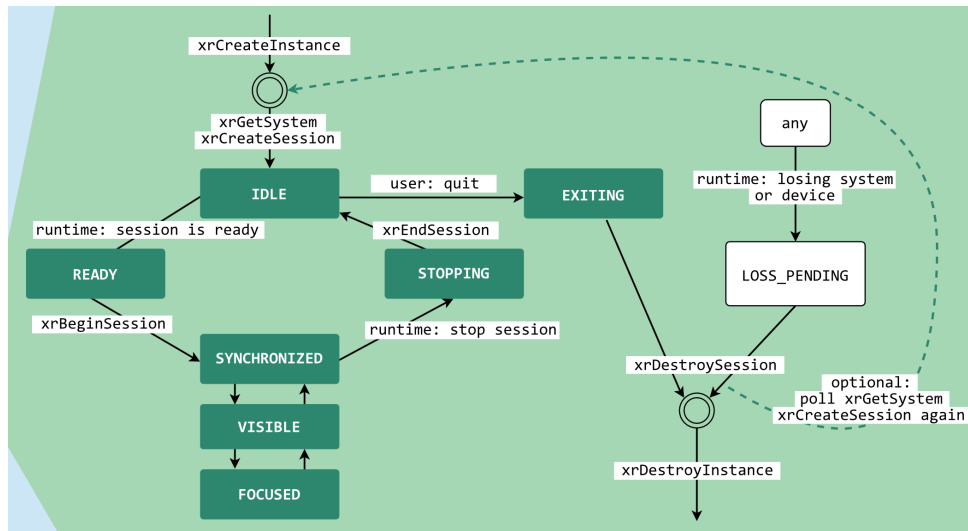


3. System information is fetched. At this point, we can use various methods to get configuration information about the view system (for instance, the “screen” size).
4. At this point, a graphic context must be created, although there is no need to have a window (unless we want to mirror what the user will see).
5. OpenXR is told which graphics API to use.
6. A “session” is created within the instance, with binding to the graphic context. Future methods will use this session.
7. The swapchain is created with specific color format, size, samples count ...
8. A reference space is created to be used in head and controllers tracking.
9. A projection layer is created for the swapchain, containing the size and offset of the rectangles associated with each eye.
10. If needed, actions are created to interact with controllers.

At this point, the session is ready. We can enter the main loop:

1. Poll events from OpenXR. Update session according to the new state (see Figure 2 for details).
2. If the session is in state `READY` , `SYNCHRONIZED` , `FOCUSED` or `VISIBLE` :
 1. Wait for the next frame and when ready, begin it.
 2. Locate the views to get the eyes positions and update the projection accordingly.
 3. Acquire the swapchain image and render to it.
 4. Release the image and end the frame.

¹³Extract of the [OpenXR Reference Guide](#)

Figure 2 — Lifecycle of an OpenXR session¹³

II - 3. MuJoCo

MuJoCo has a whole [documentation chapter](#) dedicated on visualization and rendering that is worth reading. It also contains tips for rendering to a VR headset. Other than what is explained in this chapter, there is not much knowledge of MuJoCo required to succeed in this project.

III. Implementation

Most of the OpenXR / OpenGL related code has been inspired by the [gl_example](#) pyopenxr example.

A Python file containing the whole source code is available in a [Gist on GitHub](#). You can also find a version in the Section V.

We will now see step by step how everything work.

III - 1. OpenXR initialization

The related method is `_init_xr`.

The first part of the method is about creating the OpenXR instance with the required extension. A lot of lines (from line 44 to line 63) are dedicated to make debugging work; this is not interesting. Without it, we can re-write most of the method using one single method call:

```

34 self._xr_instance = xr.create_instance(xr.InstanceCreateInfo(
35     enabled_extension_names: [xr.KHR_OPENGL_ENABLE_EXTENSION_NAME],
36     application_info=xr.ApplicationInfo(
37         application_name=APP_NAME,
38         engine_name="pyopenxr",
39         engine_version=xr.PYOPENXR_CURRENT_API_VERSION,
40         api_version=xr.Version(1, 0, xr.XR_VERSION_PATCH)
41     )
42 ))

```

`xr.KHR_OPENGL_ENABLE_EXTENSION_NAME` is a constant describing the name of the `XR_KHR_opengl_enable` extension.¹⁴ `APP_NAME` is a constant holding the name of our application that should be displayed to the user. It is defined at the beginning of the file. There is nothing really fancy here.

```

71 self._xr_system = xr.get_system(self._xr_instance,
72     xr.SystemGetInfo(xr.FormFactor.HEAD_MOUNTED_DISPLAY))
73 assert xr.enumerate_view_configurations(self._xr_instance,
74     self._xr_system)[0] == xr.ViewConfigurationType.PRIMARY_STEREO
75 views_config = xr.enumerate_view_configuration_views(self._xr_instance,
76     self._xr_system, xr.ViewConfigurationType.PRIMARY_STEREO)
77 assert len(views_config) == 2
78 assert views_config[0].recommended_image_rect_width ==
79     views_config[1].recommended_image_rect_width
80 assert views_config[0].recommended_image_rect_height ==
81     views_config[1].recommended_image_rect_height
82 self._width, self._height = views_config[0].recommended_image_rect_width,
83     views_config[0].recommended_image_rect_height
84 self._width_render = self._width * 2

```

¹⁴https://registry.khronos.org/OpenXR/specs/1.1-khr/html/xrspec.html#XR_KHR_opengl_enable

Here we do a bunch of checks to ensure everything is fine and avoid weird errors when rendering afterwards. The `enumerate_view_configuration_views` call at line 73 allows to get the image width and height, that we store. We also store an additional `_width_render` field that is simply the double of the normal width: it is the total width of our stereo render target.

The last part of this method is an ugly mixture of Python and C code to tell OpenXR to use OpenGL. We check that it contains no exception, and then we exit the method.

III - 2. OpenGL context and window

The related method is `_init_window`.

```

104 if not glfw.init():
105     raise RuntimeError("GLFW initialization failed")
106 glfw.window_hint(glfw.DOUBLEBUFFER, False)
107 glfw.window_hint(glfw.RESIZABLE, False)
108 glfw.window_hint(glfw.SAMPLES, 0) # no need for multisampling here, we
    will resolve ourselves
109 if not self._mirror_window:
110     glfw.window_hint(glfw.VISIBLE, False)
111 self._window_size = [self._width // 2, self._height // 2]
112 self._window = glfw.create_window(*self._window_size, APP_NAME, None,
    None)
113 if self._window is None:
114     raise RuntimeError("Failed to create GLFW window")
115 glfw.make_context_current(self._window)

```

We use GLFW to create a window. This window is not necessarily visible (see line 109) but even without it, an OpenGL context will be created.

The `glfw.window_hint` calls are to set the parameters of the window/context to be created. Notably, we disable double-buffering (a swapchain of two images) because we do not really care about the window rendering quality and it can save us the time of image swapping.

The window is created with half the size of one eye, so it can fit on the screen (because most VR headset have huge resolutions).

On line 114, we tell glfw to make the OpenGL context current in the current thread.

It is worth noting that we do not give any information on the OpenGL version and profile we want. This is because MuJoCo specifically requires the *Compatibility profile* (see [this MuJoCo page](#) and Section IV - 4.).

III - 3. OpenXR configuration

The related method is `_prepare_xr`.

The first part of the method has nothing really worth noting: it first creates the `graphics_binding` object that must be passed to the OpenXR session and then it creates the actual session using some C/Python code.

```

146 self._xr_swapchain = xr.create_swapchain(self._xr_session,
    xr.SwapchainCreateInfo(
147     usage_flags=xr.SWAPCHAIN_USAGE_TRANSFER_DST_BIT |
    xr.SWAPCHAIN_USAGE_COLOR_ATTACHMENT_BIT | xr.SWAPCHAIN_USAGE_SAMPLED_BIT,

```

```

148     format=GL.GL_RGBA8,
149     sample_count=1 if self._samples is None else self._samples,
150     array_size=1,
151     face_count=1,
152     mip_count=1,
153     width=self._width_render,
154     height=self._height
155 ))
    self._xr_swapchain_images =
156 xr.enumerate_swapchain_images(self._xr_swapchain,
    xr.SwapchainImageOpenGLKHR)

```

At line 145 we create the swapchain for our stereo image, hence its width being `_width_render` (which is twice one eye's image width). The format has been selected arbitrarily, the best practice would be to enumerate the available formats and select the best one from here. The usage flags¹⁵ contains:

- `TRANSFER_DST` because the swapchain image will be the destination of a pixel transfer operation (seen later)
- `COLOR_ATTACHMENT` because the swapchain image will have colored pixels on it (in most OpenXR applications this is the case)
- `SAMPLED` because the image can be multisampled

The line 155 is there to retrieve the list of images contained in the swapchain: we do it here once instead of doing it for each frame.

```

158 self._xr_projection_layer = xr.CompositionLayerProjection(
159     # Default space params are okay: identity quaternion and zero vector.
    Let's use them.
    space=xr.create_reference_space(self._xr_session,
160 xr.ReferenceSpaceCreateInfo()),
161     views = [xr.CompositionLayerProjectionView(
162         sub_image=xr.SwapchainSubImage(
163             swapchain=self._xr_swapchain,
164             image_rect=xr.Rect2Di(
165                 extent=xr.Extent2Di(self._width, self._height),
166                 offset=None if eye_index == 0 else xr.Offset2Di(x =
167 self._width) # right eye offset
168             )
169         ) for eye_index in range(2)]
170     )
171
172 self._xr_swapchain_fbo = GL.glGenFramebuffers(1)

```

Here we create the projection layer for the swapchain. It is the object that instructs the run-time where to put the rendered image in the virtual user space. It is done in 3 parts:

1. The reference space for the projection is created with the default options (the `STAGE` space type¹⁶ and default orientation and position).

¹⁵<https://registry.khronos.org/OpenXR/specs/1.1/man/html/XrSwapchainUsageFlagBits.html>

2. One view per eye is created. Both views are attached to the same swapchain (the one we created earlier) but the `image_rect` of the right eye (which defines which part of the swapchain image is displayed) is offset to the right.
3. Finally, the whole projection layer is created.

Finally at line 171, we create an empty OpenGL framebuffer for the swapchain. We will use it later.

III - 4. MuJoCo preparation

The related method is `_prepare_mujoco`.

```

178 self._mj_model = mujoco.MjModel.from_xml_path("assets/balloons.xml")
179 self._mj_data = mujoco.MjData(self._mj_model)
180 self._mj_scene = mujoco.MjvScene(self._mj_model, 1000)
181 self._mj_scene.stereo = mujoco.mjtStereo.mjSTEREO_SIDE_BY_SIDE
182
183 # We want the visualization properties set BEFORE creation of the
184 # context,
185 # otherwise we would have to call mjr_resizeOffscreen.
186 self._mj_model.vis.global_.offwidth = self._width_render
187 self._mj_model.vis.global_.offheight = self._height
188 self._mj_model.vis.quality.offsamples = 0 if self._samples is None else
189 self._samples
190
191 self._mj_context = mujoco.MjrContext(self._mj_model,
192 mujoco.mjtFontScale.mjFONTSCALE_100)
193 self._mj_camera = mujoco._structs.MjvCamera()
194 self._mj_option = mujoco.MjvOption()
195 # We do NOT want to call mjv_defaultFreeCamera
196
197 mujoco.mjv_defaultOption(self._mj_option)

```

The lines 177 and 178 are basic MuJoCo initialization. This can be done somewhere else in the code, even far sooner.

In this method, we mainly initialize the options of the MuJoCo scene and visualization objects so it can create its render context accordingly. This is done at line 190: when initializing the `MjrContext` object, it internally creates the offscreen framebuffer with the parameters we set at lines 184 to 186.

At this point, everything is ready to start the main loop.

III - 5. Frame loop - first part

The related methods are `loop`, `frame`, `_poll_xr_events` and `_wait_xr_frame`.

The main loop structure looks like this:

```

loop:
    poll_events
    if should_quit:
        stop

```

¹⁶<https://registry.khronos.org/OpenXR/specs/1.1/man/html/XrReferenceSpaceType.html>

```

if try_start_frame:
    make_a_frame

```

The `poll_events` part is made of this:

```

362 glfw.poll_events()
363 self._poll_xr_events()
364 if glfw.window_should_close(self._window):
365     self._should_quit = True

```

The `poll_events` method of glfw allows to know if the user wants to close the application on the desktop part (for instance, by closing the mirror window). We update the `_should_quit` field accordingly at line 366.

The `_poll_xr_events` method is fetching all events from the OpenXR instance and, if the event is a `SESSION_STATE_CHANGED` event, it does the following:

```

237 match self._xr_session_state:
238     case xr.SessionState.READY:
239         if not self._should_quit:
240             xr.begin_session(self._xr_session,
241                             xr.SessionBeginInfo(xr.ViewConfigurationType.PRIMARY_STEREO))
242     case xr.SessionState.STOPPING:
243         # means the session should end BUT it can start again later,
244         # this happens for instance when the user removes the headset
245         xr.end_session(self._xr_session)
246     case xr.SessionState.EXITING | xr.SessionState.LOSS_PENDING:
247         self._should_quit = True

```

If this is not clear to you, see the Session lifecycle at Figure 2.

If everything is fine and the visualization should not quit, we wait for the next XR frame:

```

210 if self._xr_session_state in [
211     xr.SessionState.READY,
212     xr.SessionState.FOCUSED,
213     xr.SessionState.SYNCHRONIZED,
214     xr.SessionState.VISIBLE,
215 ]:
216     self._xr_frame_state = xr.wait_frame(self._xr_session,
217     xr.FrameWaitInfo())
218     return True
219 return False

```

If the session is in the right state to render a frame, it waits for the frame to be ready (so we do not render faster than the device refresh rate) and *then* it returns `True`. We do *not* start the frame now: it's better to do it the moment right before we will actually render something.

III - 6. Frame loop - second part

The related methods are `frame`, `_update_mujoco`, `_update_views`, `render`, `_end_xr_frame`.

This only happens if the session is in the state to render a frame. This is how it goes:

```

371 self._update_mujoco()
372 self._update_views()
373
374 xr.begin_frame(self._xr_session, None)
375 if self._xr_frame_state.should_render:
376     self._render()
377 self._end_xr_frame()

```

The `_update_mujoco` method is really simple:

```

200 mujoco.mj_step(self._mj_model, self._mj_data)
    mujoco.mjv_updateScene(self._mj_model, self._mj_data, self._mj_option,
201 None, self._mj_camera, mujoco.mjtCatBit.mjCAT_ALL, self._mj_scene)

```

The line 199 could be done externally, it is not tied to the visualization: it only steps the physics. On the contrary, the call to `mjv_updateScene` at line 200 fetches geometries from the simulation data and stores it in the scene.

The `_update_views` method is the one that takes care of the head tracking. It goes in 3 parts: first, it fetches the `view_states` which contains, for each eye, its position, orientation and field of view. Then, it updates the projection layer accordingly and the two cameras in the MuJoCo scene to follow the eyes. Finally, it tells MuJoCo that all coordinates should be transformed in a certain way (otherwise, the world is tilted to the right).

The `_render` function is important and complex:

```

287 # We first ask to acquire a swapchain image to render onto
    image_index = xr.acquire_swapchain_image(self._xr_swapchain,
288 xr.SwapchainImageAcquireInfo())
    xr.wait_swapchain_image(self._xr_swapchain,
289 xr.SwapchainImageWaitInfo(timeout=xr.INFINITE_DURATION))
290
291 # Once we acquired it, we bind the image to our framebuffer object
292 GL.glBindFramebuffer(GL.GL_FRAMEBUFFER, self._xr_swapchain_fbo)
293 GL.glFramebufferTexture2D(
294     GL.GL_FRAMEBUFFER,
295     GL.GL_COLOR_ATTACHMENT0,
296     GL.GL_TEXTURE_2D if self._samples == None else
    GL.GL_TEXTURE_2D_MULTISAMPLE,
297     self._xr_swapchain_images[image_index].image,
298     0
299 )

```

This first part prepares the framebuffer we created at the end of Section III - 3. by attaching the current swapchain image.

`glBindFramebuffer(GL_FRAMEBUFFER, fbo)` sets the framebuffer object as the one which will receive the read and draw operations.

`glFramebufferTexture2D` attaches the image as the first color attachment of the framebuffer object.

```
302 mujoco.mjr_setBuffer(mujoco.mjtFramebuffer.mjFB_OFFSCREEN,  
    self._mj_context)  
303 mujoco.mjr_render(mujoco.MjrRect(0, 0, self._width_render, self._height),  
    self._mj_scene, self._mj_context)
```

The *real* rendering is done in the line 304. Afterwards, all is left is to copy the final image from MuJoCo's offscreen framebuffer to our own framebuffer, which has the swapchain image attached.

```
305 # We copy what MuJoCo rendered on our framebuffer object  
306 GL.glBindFramebuffer(GL.GL_READ_FRAMEBUFFER, self._mj_context.offFB0)  
307 GL.glBindFramebuffer(GL.GL_DRAW_FRAMEBUFFER, self._xr_swapchain_fbo)  
308 GL.glBlitFramebuffer(  
309     0, 0,  
310     self._width_render, self._height,  
311     0, 0,  
312     self._width_render, self._height,  
313     GL.GL_COLOR_BUFFER_BIT,  
314     GL.GL_NEAREST  
315 )
```

The first two instructions are to set which framebuffer will be read from and which one will be drawn on.

`glBlitFramebuffer` is an instruction to “copy” the pixels (the color ones in our case) from the read framebuffer to the draw one. Both framebuffers color attachments have the same size, so we put the same rectangle twice.

The rest of the method is made to downsample the rendered image and then copy it to our mirror window (if needed).

IV. Enhancements

IV - 1. Real-time simulation

For now, the code does 1 simulation step per render frame. However, due to synchronization made by OpenXR, one frame cannot be *shorter* than what it is supposed to be, so the framerate does not exceed the refresh rate of the device (for instance, 80Hz for the Oculus Rift S). This means that the simulation will update 80 times per second, no more. If the timestep set in the MuJoCo is not set to 1/80 of a second, the simulation will not be in “real-time”.

To fix that, there are two options:

- The easy one is to change the timestep of your MuJoCo model to match the refresh rate. For the Oculus Rift S, you would set the timestep to $1/80 = 0.0125s$. This however is not ideal because some simulations will not be stable at such a large timestep.
- The harder one is to change the code to do, for each frame, the amount of simulation steps needed to advance the same amount of time the frame should durate. For a frame duration of $\Delta t_{\text{frame}} = 1/f$ and a timestep of Δt_{sim} , you would advance for

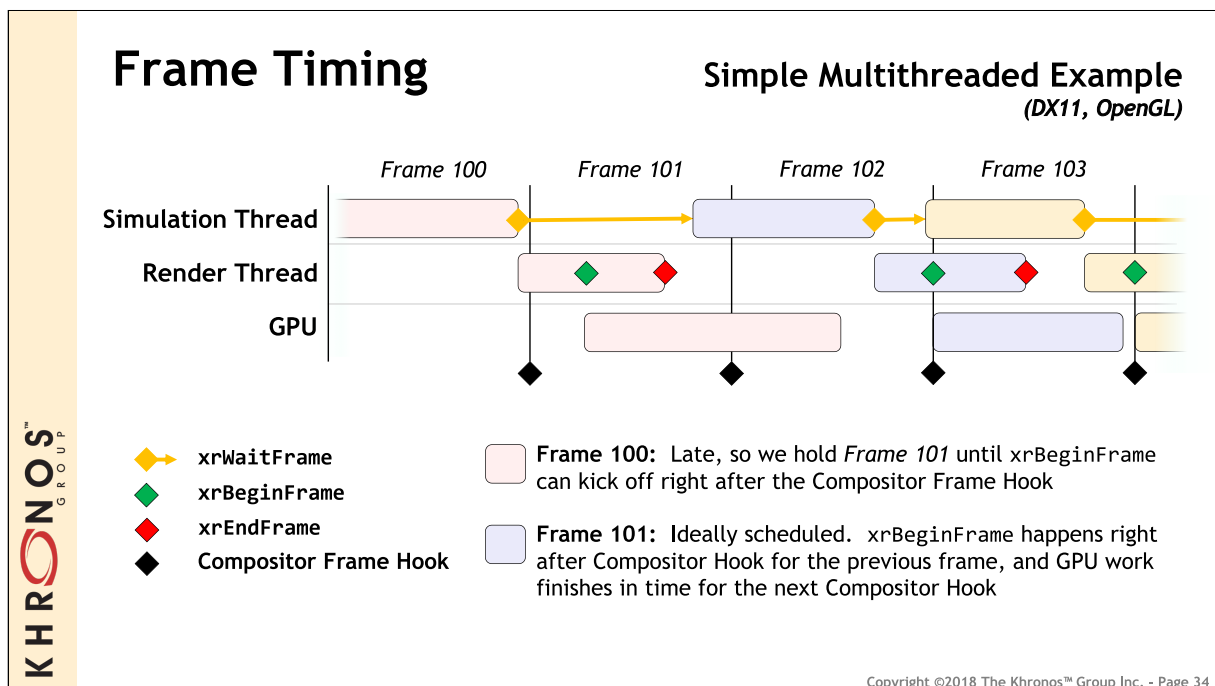
$$n_{\text{steps}} = \left\lfloor \frac{\Delta t_{\text{frame}}}{\Delta t_{\text{sim}}} \right\rfloor$$

You can get the frame duration using `_xr_frame_state.predicted_display_period` (in nanoseconds).

IV - 2. Performance

For now, everything is single-threaded. If the simulation time is long, it can reduce drastically the framerate and lead to uncomfortable VR experience. It is however possible to split our main loop in two different threads: a “simulation” thread and a “render” thread (see Figure 3).

Figure 3 — How to multithread an OpenXR application¹⁷



As you can see, the two threads will not be entirely parallel but some of their work will definitely be.

To implement this, we first need to define which work will be given to which thread:

- **Simulation thread:** event polling, frame waiting and simulation stepping.
- **Render thread:** view update, frame begin/end, rendering.

Beware: until now, we have grouped the simulation stepping and rendering preparation in one single step (`_update_mujoco`). This should now be split: `mj_step` will be kept in the simulation thread, while `mjv_updateScene` will be put in the rendering thread. This is because `mj_step` can be called while the simulation is being rendered, whereas `mjv_updateScene` cannot (because it contains the geometry data to be rendered).

Now we know where to put which work, the only thing left to do is to split the original `frame` function in two different threads, and use a synchronization structure to signal to the render thread to do its work when simulation is over (a `threading.Semaphore` is suited for this).

A race condition might happen when you remove the headset in this multithreaded setup: `xr.end_session` might be called by the simulation thread while a frame is still being rendered by the other thread. To avoid this, use another synchronization structure (a `threading.Condition` is fine).

IV - 3. Hand tracking

Hand tracking is not included in the demo file, because it is tied in how the hand is represented in the MJCF. However, here are the basic steps to implement it:

- In the MuJoCo model, add a mocap body that will receive the hand position.
- For OpenXR, we have to create an *action* that will receive data from the controller. To do that, in the Python program after the `_prepare_xr` step, add another step which follow the same steps as in [this example](#).
- To get the position and orientation in each frame, use this code:

```

1  xr.sync_actions(self._xr_session, xr.ActionsSyncInfo(active_action_sets =
2  ctypes.pointer(xr.ActiveActionSet(
3  action_set=self._action_set,
4  subaction_path=xr.NULL_PATH # wildcard to get all actions
5  )))
6  space_location = xr.locate_space(
7  space=self._action_space,
8  base_space=self._xr_projection_layer.space,
9  time=self._xr_frame_state.predicted_display_time
10 )
11 if (space_location.location_flags & xr.SPACE_LOCATION_POSITION_VALID_BIT
12     and space_location.location_flags &
13     xr.SPACE_LOCATION_ORIENTATION_VALID_BIT):
14     hand_pos = numpy.zeros(3)
15     hand_rot = numpy.zeros(4)

```

¹⁷Extract of the [OpenXR presentation at GDC 2018](#)

```

15     orientation = [space_location.pose.orientation[3],
16     *space_location.pose.orientation[3]]
17     mujoco.mjv_room2model(hand_pos, hand_rot,
    list(space_location.pose.position), orientation, self._mj_scene)
    return hand_pos, hand_rot

```

The important part here is to remember to change the quaternion format to the MuJoCo one (w, x, y, z) and to call `mjv_room2model` to automatically apply the transformations defined in `_update_views` to the pose.

IV - 4. Note on the `ContextObject` provided by *pyopenxr*

The *pyopenxr* bindings provide a pre-made class to handle most of the instance, session and swapchain work and let us focus on the interesting parts. This class is named `ContextObject`. However, it is not suitable for use in our case for two reasons:

- `ContextObject` creates one swapchain per eye, whereas we want one big swapchain containing both eyes (because this is how we want MuJoCo to render).
 - This issue could have been bypassed by creating a new swapchain ourselves and not using the premade ones. However, we loose a big part of the advantage of using this class in the first hand: simplicity.
- `ContextObject` uses an internal `OpenGLGraphics` class that handles a lot of rendering-related code. However, this class initializes the OpenGL context with the version 4.5 and *Core profile*. As we saw earlier, MuJoCo requires the *Compatibility profile*.
 - This is not bypassable without recompiling all of *pyopenxr*.

For those reasons, we made every OpenXR-related code from the ground up.

V. Annex - Source Code

```

1  import xr
2  import mujoco
3  import glfw
4  import platform
5  import ctypes
6  import numpy
7  from OpenGL import GL
8  from typing import Optional
9
10 APP_NAME = "MuJoCo XR Viewer"
11 FRUSTUM_NEAR = 0.05
12 FRUSTUM_FAR = 50
13
14 class MujocoXRViewer:
15     def __init__(self, mirror_window = False, debug = False, samples: Optional[int] = None):
16         self._mirror_window = mirror_window
17         self._debug = debug
18         self._samples = samples
19         self._should_quit = False
20
21     def __enter__(self):
22         self._init_xr()
23         self._init_window()
24         self._prepare_xr()
25         self._prepare_mujoco()
26         glfw.make_context_current(None) # To let other threads use the context if needed
27         return self
28
29     def _init_xr(self):
30         """
31         Initializes the OpenXR environment prior to session creation.
32
33         Also fetches informations about the setup, most importantly the render size.
34         """
35         extensions = [xr.KHR_OPENGL_ENABLE_EXTENSION_NAME]
36         instance_create_info = xr.InstanceCreateInfo(
37             application_info=xr.ApplicationInfo(
38                 application_name=APP_NAME,
39                 engine_name="pyopenxr",
40                 engine_version=xr.PYOPENXR_CURRENT_API_VERSION,
41                 api_version=xr.Version(1, 0, xr.XR_VERSION_PATCH)
42             )
43         )
44
45         if self._debug:
46             def debug_callback_py(severity, _type, data, _user_data):
47                 print(severity, f"{data.contents.function_name.decode()}: {data.contents.message.decode()}")
48                 return True
49
50             debug_messenger = xr.DebugUtilsMessengerCreateInfoEXT(
51                 message_severities=
52                     xr.DEBUG_UTILS_MESSAGE_SEVERITY_VERBOSE_BIT_EXT
53                     | xr.DEBUG_UTILS_MESSAGE_SEVERITY_INFO_BIT_EXT
54                     | xr.DEBUG_UTILS_MESSAGE_SEVERITY_WARNING_BIT_EXT
55                     | xr.DEBUG_UTILS_MESSAGE_SEVERITY_ERROR_BIT_EXT,
56                 message_types=
57                     xr.DEBUG_UTILS_MESSAGE_TYPE_GENERAL_BIT_EXT
58                     | xr.DEBUG_UTILS_MESSAGE_TYPE_VALIDATION_BIT_EXT
59                     | xr.DEBUG_UTILS_MESSAGE_TYPE_PERFORMANCE_BIT_EXT
60                     | xr.DEBUG_UTILS_MESSAGE_TYPE_CONFORMANCE_BIT_EXT,
61                 user_callback=xr.PFN_xrDebugUtilsMessengerCallbackEXT(debug_callback_py)
62             )
63             instance_create_info.next = ctypes.cast(ctypes.pointer(debug_messenger), ctypes.c_void_p)
64             extensions.append(xr.EXT_DEBUG_UTILS_EXTENSION_NAME)
65
66         instance_create_info.enabled_extension_names = extensions
67         self._xr_instance = xr.create_instance(instance_create_info)
68
69         # The following fetches important informations about the setup
70         # (mainly rendering size)
71         self._xr_system = xr.get_system(self._xr_instance, xr.SystemGetInfo(xr.FormFactor.HEAD_MOUNTED_DISPLAY))
72         assert xr.enumerate_view_configurations(self._xr_instance, self._xr_system)[0] ==
73         xr.ViewConfigurationType.PRIMARY_STEREO

```



```

74     views_config = xr.enumerate_view_configuration_views(self._xr_instance, self._xr_system,
75     xr.ViewConfigurationType.PRIMARY_STEREO)
76     assert len(views_config) == 2
77     assert views_config[0].recommended_image_rect_width == views_config[1].recommended_image_rect_width
78     assert views_config[0].recommended_image_rect_height == views_config[1].recommended_image_rect_height
79
80     self._width, self._height = views_config[0].recommended_image_rect_width,
81     views_config[0].recommended_image_rect_height
82     self._width_render = self._width * 2
83
84     pxrGetOpenGLGraphicsRequirementsKHR = ctypes.cast(
85         xr.get_instance_proc_addr(
86             self._xr_instance,
87             "xrGetOpenGLGraphicsRequirementsKHR",
88         ),
89         xr.PFN_xrGetOpenGLGraphicsRequirementsKHR
90     )
91     graphics_result = pxrGetOpenGLGraphicsRequirementsKHR(
92         self._xr_instance,
93         self._xr_system,
94         ctypes.byref(xr.GraphicsRequirementsOpenGLKHR())
95     )
96     graphics_result = xr.exception.check_result(xr.Result(graphics_result))
97     if graphics_result.is_exception():
98         raise graphics_result
99
100 def _init_window(self):
101     """
102     Initializes the GLFW window (and make it hidden if mirrored mode is disabled).
103
104     Creates the OpenGL context that will be used.
105     """
106     if not glfw.init():
107         raise RuntimeError("GLFW initialization failed")
108     glfw.window_hint(glfw.DOUBLEBUFFER, False)
109     glfw.window_hint(glfw.RESIZABLE, False)
110     glfw.window_hint(glfw.SAMPLES, 0) # no need for multisampling here, we will resolve ourselves
111     if not self._mirror_window:
112         glfw.window_hint(glfw.VISIBLE, False)
113     self._window_size = [self._width // 2, self._height // 2]
114     self._window = glfw.create_window(*self._window_size, APP_NAME, None, None)
115     if self._window is None:
116         raise RuntimeError("Failed to create GLFW window")
117     glfw.make_context_current(self._window)
118     # Attempt to disable vsync on the desktop window or
119     # it will interfere with the OpenXR frame loop timing
120     glfw.swap_interval(0)
121
122 def _prepare_xr(self):
123     """
124     Creates the OpenXR session and prepares everything to launch the frames loop.
125     """
126     if platform.system() == 'Windows':
127         from OpenGL import WGL
128         graphics_binding = xr.GraphicsBindingOpenGLWin32KHR()
129         graphics_binding.h_dc = WGL.wglGetCurrentDC()
130         graphics_binding.h_glr = WGL.wglGetCurrentContext()
131     else:
132         from OpenGL import GLX
133         graphics_binding = xr.GraphicsBindingOpenGLXlibKHR()
134         graphics_binding.x_display = GLX.glXGetCurrentDisplay()
135         graphics_binding.glx_context = GLX.glXGetCurrentContext()
136         graphics_binding.glx_drawable = GLX.glXGetCurrentDrawable()
137
138     self._xr_session = xr.create_session(
139         self._xr_instance,
140         xr.SessionCreateInfo(
141             0,
142             self._xr_system,
143             next=ctypes.cast(ctypes.pointer(graphics_binding), ctypes.c_void_p)
144         )
145     )
146     self._xr_session_state = xr.SessionState.IDLE
147
148     self._xr_swapchain = xr.create_swapchain(self._xr_session, xr.SwapchainCreateInfo(
149         usage_flags=xr.SWAPCHAIN_USAGE_TRANSFER_DST_BIT | xr.SWAPCHAIN_USAGE_COLOR_ATTACHMENT_BIT |
150         xr.SWAPCHAIN_USAGE_SAMPLED_BIT,

```

```

148         format=GL.GL_RGBA8,
149         sample_count=1 if self._samples is None else self._samples,
150         array_size=1,
151         face_count=1,
152         mip_count=1,
153         width=self._width_render,
154         height=self._height
155     ))
156     self._xr_swapchain_images = xr.enumerate_swapchain_images(self._xr_swapchain, xr.SwapchainImageOpenGLKHR)
157
158     self._xr_projection_layer = xr.CompositionLayerProjection(
159         # Default space params are okay: identity quaternion and zero vector. Let's use them.
160         space=xr.create_reference_space(self._xr_session, xr.ReferenceSpaceCreateInfo()),
161         views = [xr.CompositionLayerProjectionView(
162             sub_image=xr.SwapchainSubImage(
163                 swapchain=self._xr_swapchain,
164                 image_rect=xr.Rect2Di(
165                     extent=xr.Extent2Di(self._width, self._height),
166                     offset=None if eye_index == 0 else xr.Offset2Di(x = self._width) # right eye offset
167                 )
168             )
169         ) for eye_index in range(2)]
170     )
171
172     self._xr_swapchain_fbo = GL.glGenFramebuffers(1)
173
174     def _prepare_mujoco(self):
175         """
176         Prepares the MuJoCo environment.
177         """
178         self._mj_model = mujoco.MjModel.from_xml_path("assets/balloons.xml")
179         self._mj_data = mujoco.MjData(self._mj_model)
180         self._mj_scene = mujoco.MjvScene(self._mj_model, 1000)
181         self._mj_scene.stereo = mujoco.mjtStereo.mjSTEREO_SIDE_BY_SIDE
182
183         # We want the visualization properties set BEFORE creation of the context,
184         # otherwise we would have to call mjr_resizeOffscreen.
185         self._mj_model.vis.global_.offwidth = self._width_render
186         self._mj_model.vis.global_.offheight = self._height
187         self._mj_model.vis.quality.offsamples = 0 if self._samples is None else self._samples
188
189         self._mj_context = mujoco.MjrContext(self._mj_model, mujoco.mjtFontScale.mjFONTSCALE_100)
190         self._mj_camera = mujoco._structs.MjvCamera()
191         self._mj_option = mujoco.MjvOption()
192         # We do NOT want to call mjr_defaultFreeCamera
193
194         mujoco.mjv_defaultOption(self._mj_option)
195
196     def _update_mujoco(self):
197         """
198         Updates MuJoCo for one frame.
199         """
200         mujoco.mj_step(self._mj_model, self._mj_data)
201         mujoco.mjv_updateScene(self._mj_model, self._mj_data, self._mj_option, None, self._mj_camera,
202                                mujoco.mjtCatBit.mjCAT_ALL, self._mj_scene)
203
204     def _wait_xr_frame(self):
205         """
206         Wait to begin the next OpenXR frame.
207
208         Returns:
209             bool: whether or not we should update the scene and maybe render it.
210         """
211         if self._xr_session_state in [
212             xr.SessionState.READY,
213             xr.SessionState.FOCUSED,
214             xr.SessionState.SYNCHRONIZED,
215             xr.SessionState.VISIBLE,
216         ]:
217             self._xr_frame_state = xr.wait_frame(self._xr_session, xr.FrameWaitInfo())
218             return True
219         return False
220
221     def _end_xr_frame(self):
222         xr.end_frame(self._xr_session, xr.FrameEndInfo(
223             self._xr_frame_state.predicted_display_time,
224             xr.EnvironmentBlendMode.OPAQUE,
225             layers=[ctypes.byref(self._xr_projection_layer)] if self._xr_frame_state.should_render else []
226         ))

```

```

226
227 def _poll_xr_events(self):
228     while True:
229         try:
230             event_buffer = xr.poll_event(self._xr_instance)
231             event_type = xr.StructureType(event_buffer.type)
232             if event_type == xr.StructureType.EVENT_DATA_SESSION_STATE_CHANGED:
233                 event = ctypes.cast(
234                     ctypes.byref(event_buffer),
235                     ctypes.POINTER(xr.EventDataSessionStateChanged)).contents
236                 self._xr_session_state = xr.SessionState(event.state)
237                 match self._xr_session_state:
238                     case xr.SessionState.READY:
239                         if not self._should_quit:
240                             xr.begin_session(self._xr_session,
241 xr.SessionBeginInfo(xr.ViewConfigurationType.PRIMARY_STEREO))
242                     case xr.SessionState.STOPPING:
243                         # means the session should end BUT it can start again later,
244                         # this happens for instance when the user removes the headset
245                         xr.end_session(self._xr_session)
246                     case xr.SessionState.EXITING | xr.SessionState.LOSS_PENDING:
247                         self._should_quit = True
248             except xr.EventUnavailable:
249                 break # We got all events
250
251 def _update_views(self):
252     _, view_states = xr.locate_views(self._xr_session, xr.ViewLocateInfo(
253         xr.ViewConfigurationType.PRIMARY_STEREO,
254         self._xr_frame_state.predicted_display_time,
255         self._xr_projection_layer.space,
256     ))
257     for eye_index, view_state in enumerate(view_states):
258         self._xr_projection_layer.views[eye_index].fov = view_state.fov
259         self._xr_projection_layer.views[eye_index].pose = view_state.pose
260
261         cam = self._mj_scene.camera[eye_index]
262         cam.pos = list(view_state.pose.position)
263         cam.frustum_near = FRUSTUM_NEAR
264         cam.frustum_far = FRUSTUM_FAR
265         cam.frustum_bottom = numpy.tan(view_state.fov.angle_down) * FRUSTUM_NEAR
266         cam.frustum_top = numpy.tan(view_state.fov.angle_up) * FRUSTUM_NEAR
267         cam.frustum_center = 0.5 * (numpy.tan(view_state.fov.angle_left) + numpy.tan(view_state.fov.angle_right)) *
FRUSTUM_NEAR
268         # no need to set left/right as it will be computed using center
269
270         rot_quat = list(view_state.pose.orientation)
271         # Guess what? OpenXR quaternions are in form (x, y, z, w)
272         # while MuJoCo quaternions are in form (w, x, y, z)...
273         rot_quat = [rot_quat[3], *rot_quat[0:3]]
274
275         forward, up = numpy.zeros(3), numpy.zeros(3)
276         mujoco.mju_rotVecQuat(forward, [0, 0, -1], rot_quat)
277         mujoco.mju_rotVecQuat(up, [0, 1, 0], rot_quat)
278         cam.forward, cam.up = forward.tolist(), up.tolist()
279
280         self._mj_scene.enabletransform = True
281         self._mj_scene.rotate[0] = numpy.cos(0.25 * numpy.pi)
282         self._mj_scene.rotate[1] = numpy.sin(-0.25 * numpy.pi)
283
284 def _render(self):
285     """
286     Renders the scene in the swapchain and eventually mirrors it on the window if needed.
287     """
288     # We first ask to acquire a swapchain image to render onto
289     image_index = xr.acquire_swapchain_image(self._xr_swapchain, xr.SwapchainImageAcquireInfo())
290     xr.wait_swapchain_image(self._xr_swapchain, xr.SwapchainImageWaitInfo(timeout=xr.INFINITE_DURATION))
291
292     # Once we acquired it, we bind the image to our framebuffer object
293     GL.glBindFramebuffer(GL.GL_FRAMEBUFFER, self._xr_swapchain_fbo)
294     GL.glFramebufferTexture2D(
295         GL.GL_FRAMEBUFFER,
296         GL.GL_COLOR_ATTACHMENT0,
297         GL.GL_TEXTURE_2D if self._samples == None else GL.GL_TEXTURE_2D_MULTISAMPLE,
298         self._xr_swapchain_images[image_index].image,
299         0
300     )
301
302     # We ask MuJoCo to render on its own offscreen framebuffer

```

```

302 mujoco.mjv_setBuffer(mujoco.mjtFramebuffer.mjFB_OFFSCREEN, self._mj_context)
303 mujoco.mjv_render(mujoco.MjrRect(0, 0, self._width_render, self._height), self._mj_scene, self._mj_context)
304
305 # We copy what MuJoCo rendered on our framebuffer object
306 GL.glBindFramebuffer(GL.GL_READ_FRAMEBUFFER, self._mj_context.offFBO)
307 GL.glBindFramebuffer(GL.GL_DRAW_FRAMEBUFFER, self._xr_swapchain_fbo)
308 GL.glBlitFramebuffer(
309     0, 0,
310     self._width_render, self._height,
311     0, 0,
312     self._width_render, self._height,
313     GL.GL_COLOR_BUFFER_BIT,
314     GL.GL_NEAREST
315 )
316
317 if self._mirror_window:
318     # We copy the data from the MuJoCo buffer to the window one (0 is the default window fbo)
319     if self._samples is not None:
320         # We first resolve multi-sample if needed
321         GL.glBindFramebuffer(GL.GL_DRAW_FRAMEBUFFER, self._mj_context.offFBO_r)
322         GL.glBlitFramebuffer(
323             0, 0,
324             self._width_render, self._height,
325             0, 0,
326             self._width_render, self._height,
327             GL.GL_COLOR_BUFFER_BIT,
328             GL.GL_NEAREST
329         )
330         GL.glBindFramebuffer(GL.GL_READ_FRAMEBUFFER, self._mj_context.offFBO_r)
331
332     # We then copy the data to the window
333     GL.glBindFramebuffer(GL.GL_DRAW_FRAMEBUFFER, 0)
334     GL.glBlitFramebuffer(
335         0, 0,
336         self._width, self._height, # one eye only (left)
337         0, 0,
338         *self._window_size,
339         GL.GL_COLOR_BUFFER_BIT,
340         0x90BA # EXT_framebuffer_multisample_blit_scaled, SCALED_RESOLVE_FASTEST_EXT
341     )
342     xr.release_swapchain_image(self._xr_swapchain, xr.SwapchainImageReleaseInfo())
343
344 def __exit__(self, exc_type, exc_value, traceback):
345     if self._window is not None:
346         glfw.make_context_current(self._window)
347         if self._xr_swapchain_fbo is not None:
348             GL.glDeleteFramebuffers(1, [self._xr_swapchain_fbo])
349             self._xr_swapchain_fbo = None
350         glfw.terminate()
351     if self._xr_swapchain is not None:
352         xr.destroy_swapchain(self._xr_swapchain)
353     if self._xr_session is not None:
354         xr.destroy_session(self._xr_session)
355     if self._xr_instance is not None:
356         # # may break on Linux SteamVR
357         # xr.destroy_instance(self._xr_instance)
358         pass # does not seem to work
359     glfw.terminate()
360
361 def frame(self):
362     glfw.poll_events()
363     self._poll_xr_events()
364     if glfw.window_should_close(self._window):
365         self._should_quit = True
366
367     if self._should_quit:
368         return
369
370     if self._wait_xr_frame():
371         self._update_mujoco()
372         self._update_views()
373
374     xr.begin_frame(self._xr_session, None)
375     if self._xr_frame_state.should_render:
376         self._render()
377     self._end_xr_frame()
378
379 def loop(self):

```

```
380     glfw.make_context_current(self._window)
381     while not self._should_quit:
```